


# Matrix multiplication

- $A$  is a  $n \times l$  matrix with elements  $a_{ij}$   
 $B$  is a  $l \times m$  matrix with elements  $b_{ij}$   
 $C$  is a  $n \times m$  matrix with elements  $c_{ij}$ 
  - the product  $C = A \times B$  is computed as the dot product of row  $i$  in  $A$  with column  $j$  in  $B$

$$c_{i,j} = \sum_{k=0}^{l-1} a_{ik} b_{kj}$$



# Sequential matrix multiplication

- The sequential code to multiply two matrices  $A$  (of size  $n \times l$ ) and  $B$  (of size  $l \times m$ ) is

```
for (i=0; i<n; i++) {  
    for (j=0; j<m; j++) {  
        C[i][j] = 0.0;  
        for (k=0; k<l; k++) {  
            C[i][j] += A[i][k]*B[k][j];  
        }  
    }  
}
```

- Time complexity is  $O(n^3)$ 
  - ♦ for each element in  $C$  we do  $l$  multiplications and additions

# Memory access in matrix multiplication

- A straight forward implementation of matrix multiplication is very slow because it accesses memory inefficiently

- accesses to matrices  $A$  and  $C$  are efficient, but accesses to  $B$  are done with a stride equal to the row length  $m$

```
for (i=0; i<n; i++)  
    for (j=0; j<m; j++) {  
        C[i][j] = 0.0;  
        for (k=0; k<l; k++)  
            C[i][j] += A[i][k]*B[k][j];  
    }
```

- In C and C++, elements in a row are stored in consecutive memory locations
  - elements in a column are located far from each other in memory
  - we get a cache miss for every access of the  $B$  matrix
- Leads to inefficient cache memory utilization
  - the processor spends most of its time waiting for memory accesses

# Memory allocation in C

- Matrices should be allocated as a consecutive block of memory
  - the whole matrix can then be sent in a single message without first copying it to a contiguous message buffer
  - accesses are more efficient, because the elements are located in consecutive memory positions
  - less cache misses, better use of automatic prefetching
- Modern processors use prefetching
  - when a regular memory access pattern is detected, the next cache line is automatically brought in to main memory before the data is actually needed
  - cache lines are typically of length 64 bytes
- Should arrange memory accesses to take advantage of the cache memory and the prefetch mechanism

# Static memory allocation

## ■ Static memory allocation

- the elements in the matrix are stored in memory as a contiguous block in row-major order

```
const int N = 1000;  
double X[N][N];  
... use the matrix X
```

## ■ We can also read in the value of $N$

- the matrix must be declared in a scope where  $N$  is initialized

## ■ Static memory is allocated on the stack

- there is an upper limit on how large blocks of memory can be allocated on the stack

```
int N;  
printf("Give N? ");  
scanf("%d", &N);  
  
double X[N][N];  
for (i=0; i<N; i++) {  
    for (j=0; j<N; j++) {  
        X[i][j] = i+j;  
    }  
}
```

# Dynamic memory allocation

- In C, memory is dynamically allocated with the *malloc* (or *calloc*) system function

```
void * malloc(size_t SIZE)
void * calloc(size_t COUNT, size_t ELTSIZE)
```

  - *calloc* initializes each element to zero
- In C++ memory is dynamically allocated with *new*
  - Example: `int *v = new int[size];`
- Dynamic memory is allocated on the heap
  - there are no limits on the size of memory blocks, except the amount of memory available in the system
- Dynamic memory allocation can be a slow procedure
  - should not be called inside the innermost loops

# Allocating a matrix

- Allocation as a one-dimensional array of size *rows\*cols*

```
double *M;  
M = (double *) malloc(rows*cols*sizeof(double));  
/* Set matrix M to zero */  
for (i=0; i<rows; i++)  
    for (j=0; j<cols; j++)  
        M[i*cols+j] = 0.0;
```

- Have to calculate the address expressions explicitly in the code
  - can also use a macro definition to do the address calculation

```
#define MAT(i,j) (M[i*cols+j]);  
.  
.  
.  
for (i=0; i<rows; i++)  
    for (j=0; j<cols; j++)  
        MAT(i,j) = 0.0;
```


# Allocating a 2D matrix

- Allocation as a two-dimensional array, one row at a time

```
double **M;  
M = (double **) malloc(rows*sizeof(double *));  
for (i=0; i<rows; i++)  
    M[i] = (double *) malloc(cols*sizeof(double));  
.  
.  
.  
for (i=0; i<rows; i++)  
    for (j=0; j<cols; j++)  
        M[i][j] = 0.0;
```

- Gives a non-contiguous allocation

- each row is separately allocated and can be placed anywhere in memory
- can not send the matrix to another process without copying it to a contiguous message buffer






# Allocating a contiguous 2D matrix

## ■ Alternative method using a contiguous block of memory

```
double **M;    /* Row pointers */
double *Mb;    /* Where data will be stored */
M = (double **) malloc(rows*sizeof(double *));
Mb = (double *) malloc(rows*cols*sizeof(double));
/* Initialize pointers to rows in the matrix */
for (i=0; i<rows; i++)
    M[i] = Mb + i*rows
. . .
for (i=0; i<rows; i++)
    for (j=0; j<cols; j++)
        M[i][j] = 0.0;
```



# Parallel matrix multiplication

- We present a parallel algorithm for matrix multiplication based on block decomposition
  - each process computes a rectangular block (a sub-matrix) of the result
  - suitable for an implementation on distributed memory
- All the elements in the result matrix  $c_{i,j}$  can be computed in parallel
  - to compute  $c_{i,j}$  we only need to read row  $i$  of  $A$  and column  $j$  of  $B$
- There are also methods based on row decomposition
  - each process computes a number of rows of the result matrix
  - suitable for a shared memory implementation, since all processes need access to the whole  $B$  matrix

# Parallel matrix multiplication with Fox's algorithm

- For simplicity we assume that
  - the matrices are square and of order  $n$  (i.e., they are  $n \times n$  matrices)
  - the number of processes is  $n^2$
- The processes are arranged in a 2-dimensional grid
- Each process is assigned one element of the matrix
  - process  $(i,j)$  (= process with rank  $i*n+j$ ) has elements  $a_{ij}$ ,  $b_{ij}$  and  $c_{ij}$
  - we will later in the agglomeration stage modify the algorithm so that each process operates on a square block of elements (a submatrix)
- We assume that the data is already distributed among the processes

# Fox's algorithm

## ■ The algorithm proceeds in $n$ stages

- one stage for each term in the dot product of row  $i$  and column  $j$   
$$c_{ij} = a_{i0} * b_{0j} + a_{i1} * b_{1j} + \dots + a_{i,n-1} * b_{n-1,j}$$

## ■ Algorithm for process $(i,j)$

- **Stage 0:**  $c_{ij} = a_{ij} * b_{ij}$ 
  - multiply the diagonal entry of  $A$  in the own row by the own element of  $B$
- **Stage 1:**  $c_{ij} += a_{i,i+1} * b_{i+1,j}$ 
  - multiply the element one step to the right of the diagonal (in the own row) in  $A$  by the element one step below the own element of  $B$
- ...
- **Stage  $k$ :**  $c_{ij} += a_{i,i+k} * b_{i+k,j}$ 
  - multiply the element  $k$  columns to the right of the diagonal of  $A$  by the element  $k$  rows below the own element of  $B$

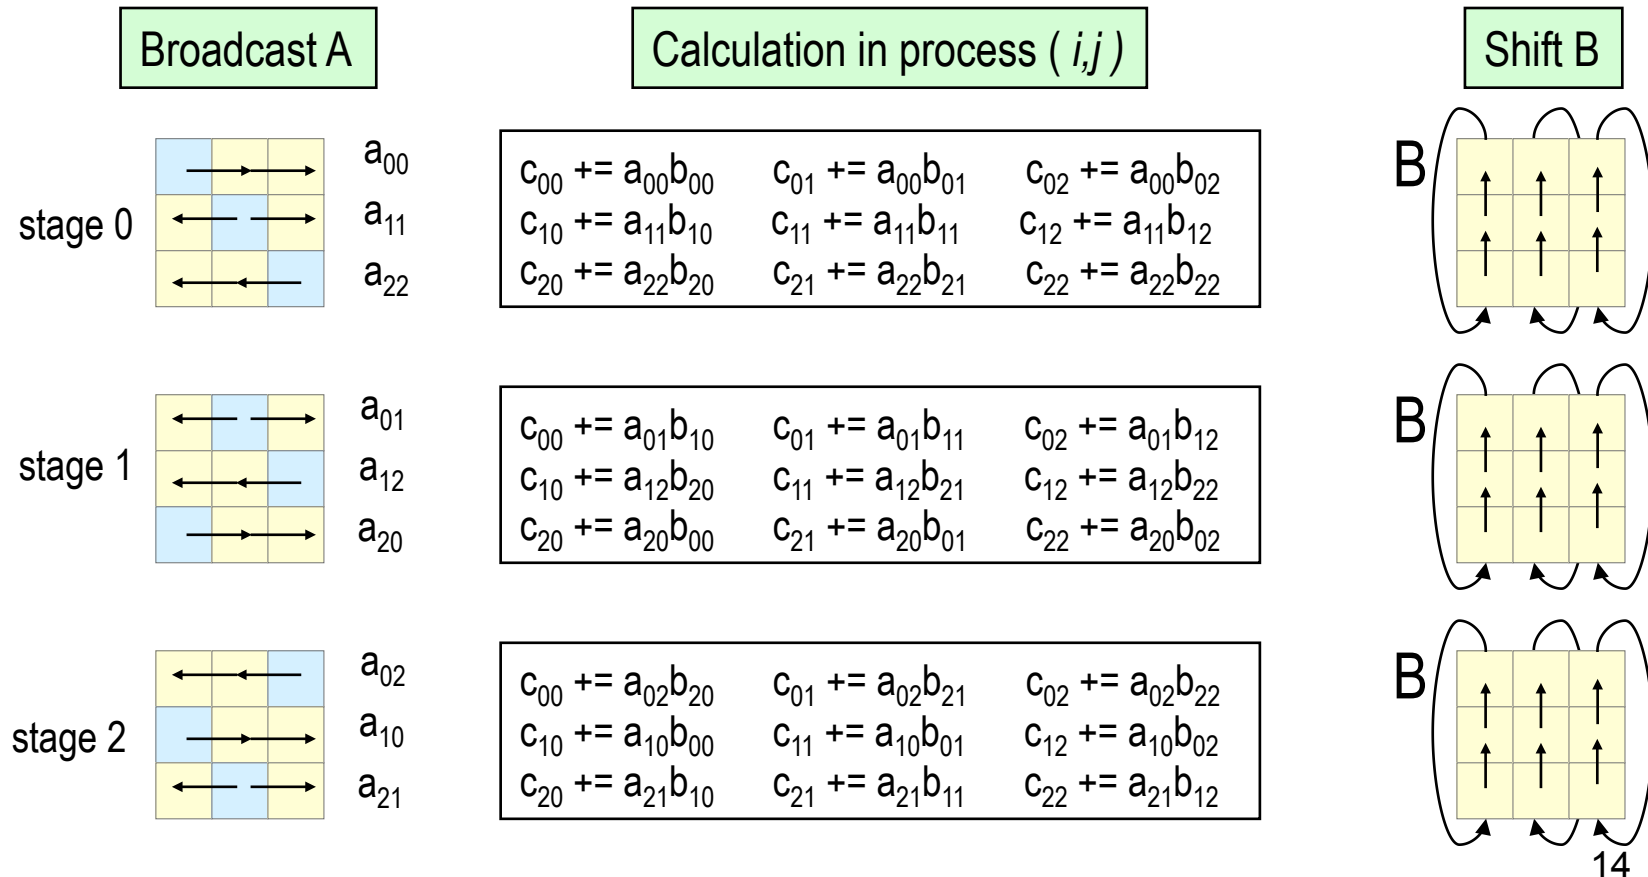
## Fox's algorithm (cont.)

- Row and column subscripts are calculated *modulo n*
  - **Stage  $k$ :**  $k' = (i+k) \bmod n$ ;  $c_{ij} += a_{i,k'} * b_{k',j}$
- The dot product  $c_{ij}$  will be computed in the order
  - $a_{ii} * b_{ij} + a_{i,i+1} * b_{i+1,j} + \dots + a_{i,n-1} * b_{n-1,j} + a_{i0} * b_{0j} + \dots + a_{i,i-1} * b_{i-1,j}$
- Each process has to get the elements  $a_{i,k'}$  and  $b_{k',j}$  from the other processes in row  $i$  and column  $j$ 
  - broadcast  $a_{i,k'}$  to all processes in row  $i$
  - calculate  $a_{i,k'} * b_{k',j}$
  - shift the value of  $b$  one step up in column  $j$  (cyclically)

# Illustration


## ■ Stage $k$ :

- broadcast the element of  $A$   $k'$  steps to the right of the diagonal to all processes in the same row
- multiply the received element of  $A$  with the current element of  $B$
- shift the element of  $B$  up in the same column (circularly)



# Block decomposition

- Instead of giving one process one element we give each process a square block of elements
  - the matrices are of size  $n \times n$
  - we use a square grid of  $p$  processes where the number of rows and columns,  $\sqrt{p}$ , evenly divides  $n$
- Each process stores a submatrix of size  $n' \times n'$ , where  $n' = n/\sqrt{p}$
- Example:
  - 16 x 16 matrix ( $n=16$ )
  - 4 processes arranged as a 2x2 grid ( $p = 4$ )
  - each process stores a 8 x 8 submatrix ( $n' = 8$ )



# Submatrices


- Let the matrix  $A_{ij}$  be the  $n' * n'$  submatrix of  $A$  whose first entry is  $a_{i*n', j*n'}$ 
  - process  $(i,j)$  is assigned submatrices  $A_{ij}$ ,  $B_{ij}$  and  $C_{ij}$
  - submatrices can be multiplied as scalar elements, just use matrix multiplication instead of scalar multiplication
- Example:
  - 4x4 matrix divided among 4 processes
  - $n = p = 4$ ,  $n' = 4 / \sqrt{4} = 2$

$$A_{0,0} = \begin{pmatrix} a_{00} & a_{01} \\ a_{10} & a_{11} \end{pmatrix}$$

$$A_{0,1} = \begin{pmatrix} a_{02} & a_{03} \\ a_{12} & a_{13} \end{pmatrix}$$

$$A_{1,0} = \begin{pmatrix} a_{20} & a_{21} \\ a_{30} & a_{31} \end{pmatrix}$$

$$A_{1,1} = \begin{pmatrix} a_{22} & a_{23} \\ a_{32} & a_{33} \end{pmatrix}$$






# Fox's algorithm with block decomposition

- $A_{ij}$ ,  $B_{ij}$  and  $C_{ij}$  denotes submatrices as defined on previous slide
- Algorithm for process  $(i,j)$ :

```
q = sqrt(p);  
dest = ((i-1) mod q, j); /* Process one step up */  
source = (i+1) mod q, j); /* Process one step below */  
  
for (stage=0; stage<q; stage++) {  
    k_prime=(i+stage) mod q;  
    Broadcast A[i,k_prime] to the processes in row i;  
    C[i,j] += A[i,k_prime]*B[k_prime,j]; /* Multiply */  
    Send B[k_prime,j] to dest;  
    Receive B[(k_prime+1) mod q,j] from source;  
}
```

# Example: 3x3 processes



- Broadcast diagonal blocks of  $A$  to processes in the same row
  - Multiply submatrices of  $A$  and  $B$  into  $C$
  - Shift blocks of  $B$  cyclically up one step
- 
- Broadcast blocks one step to the right of the diagonal in  $A$  to processes in the same row
  - Multiply submatrices  $A$  and  $B$  into  $C$
  - Shift blocks of  $B$  cyclically up one step
- 
- Broadcast blocks two steps to the right of the diagonal in  $A$  to processes in the same row
  - Multiply submatrices  $A$  and  $B$  into  $C$
  - Shift blocks of  $B$  cyclically up one step

# Code for Fox's algorithm

```
dest = (my_row+q-1)%q;    /* Destination for circular shift in columns */
source = (my_row+1)%q;    /* Source for circular shift in columns */

/* Allocate storage for temporary local matrix */
tmp = (float *) malloc(sizeof(float)*N_local*N_local);

settozero(C_local, N_local); /* Set the result matrix to zero */

for (stage=0; stage<q; stage++) {
    bcast_root = (my_row+stage)%q;    /* Process that does the broadcast */
    if (bcast_root == my_col) {
        /* Send to all other processes in the same row */
        MPI_Bcast(A_local, N_local*N_local, MPI_FLOAT, bcast_root, row_comm);
        /* Multiply the submatrices */
        matrixmult(A_local, B_local, C_local, N_local);
    } else {
        /* Receive submatrix of A from the process that broadcasts */
        MPI_Bcast(tmp, N_local*N_local, MPI_FLOAT, bcast_root, row_comm);
        /* Multiply it with own submatrix B_local */
        matrixmult(tmp, B_local, C_local, N_local);
    }
    /* Send submatrix of B up and receive a new from below */
    MPI_Sendrecv_replace(B_local, N_local*N_local, MPI_FLOAT, dest,
                        datatag, source, datatag, col_comm, &status);
}
```

# Implementing Fox's algorithm

- The implementation will contain the following steps
  - initialise MPI
  - check that we have a square number of processes (4, 9, 16, 25, ...)
  - read the size of the matrices  $N$
  - check that the number of elements in the matrices is evenly divisible by the square root of the number of processes
  - allocate memory for the input matrices and the local sub-matrices
  - read in the input matrices  $A$  and  $B$  from files
  - create a 2-dimensional process grid
  - create communicators for rows and columns in the process grid
  - distribute the matrices  $A$  and  $B$  to the processes so that each process gets its own submatrices  $A_{local}$  and  $B_{local}$
  - do the matrix multiplication with Fox's algorithm
  - collect the submatrices  $C_{local}$  from each process into a result-matrix  $C$
  - write the result to a file
  - compare the result with a known correct result (available in a file)